

NA 565 - Fall 2024

Linear Algebra Review

August 26, 2024



- ▶ Linear algebra is the workhorse of modern science and engineering.
- ▶ Digital systems and modern computations use algebraic or time-difference (discrete) equations for implementation.
- ▶ Most successful algorithms in practice boils down to doing linear algebra (machine learning, optimization, control, estimation, computer vision, etc.).
- ▶ A secret to share: If you know least squares and can solve problems, you could claim to be an engineer (unless you're asked to show proof!).

- ▶ The identity matrix is a square matrix denoted I that has ones down the main diagonal and zeroes elsewhere.
- ▶ Here are some examples of 1×1 , 2×2 , 3×3 , and 4×4 identity matrices.

$$[1], \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

- ▶ The notation I_n means an $n \times n$ identity matrix.

We identify \mathbb{R}^n with the set of all n -column vectors with real entries

$$\{(x_1, x_2, \dots, x_n) \mid x_i \in \mathbb{R}, 1 \leq i \leq n\}$$
$$\iff \mathbb{R}^n \iff \left\{ \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \mid x_i \in \mathbb{R}, 1 \leq i \leq n \right\}$$

For all real numbers α and β , and all vectors x and y in \mathbb{R}^n we have scalar multiplication and vector addition:

$$\begin{aligned}\alpha x + \beta y &= \alpha \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} + \beta \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \\ &= \begin{bmatrix} \alpha x_1 \\ \alpha x_2 \\ \vdots \\ \alpha x_n \end{bmatrix} + \begin{bmatrix} \beta y_1 \\ \beta y_2 \\ \vdots \\ \beta y_n \end{bmatrix} = \begin{bmatrix} \alpha x_1 + \beta y_1 \\ \alpha x_2 + \beta y_2 \\ \vdots \\ \alpha x_n + \beta y_n \end{bmatrix} \in \mathbb{R}^n.\end{aligned}$$

Euclidean Norm or “Length” of a Vector

Definition

Let $v = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}$ be a vector in \mathbb{R}^n . The Euclidean norm of v , denoted $\|v\|$, is defined as

$$\|v\| := \sqrt{(v_1)^2 + (v_2)^2 + \cdots + (v_n)^2} = \sqrt{\sum_{i=1}^n (v_i)^2} = \sqrt{v^T v}$$

Properties of the Norm of a vector

All norms satisfy the following properties

1 For all vectors $v \in \mathbb{R}^n$, $\|v\| \geq 0$ and moreover,
 $\|v\| = 0 \iff v = 0$.

2 For any real number $\alpha \in \mathbb{R}$ and vector $v \in \mathbb{R}^n$,

$$\|\alpha v\| = |\alpha| \cdot \|v\|.$$

3 For any pair of vectors v and w in \mathbb{R}^n ,

$$\|v + w\| \leq \|v\| + \|w\|.$$

A vector $v \in \mathbb{R}^n$ is a *linear combination* of $\{u_1, u_2, \dots, u_m\} \subset \mathbb{R}^n$ if there exist real numbers $\alpha_1, \alpha_2, \dots, \alpha_m$ such that

$$v = \alpha_1 u_1 + \alpha_2 u_2 + \cdots + \alpha_m u_m.$$

Remark

$A \subset B$ means A is a subset of B . This means that B includes or contains A . For example, $\{1,2\} \subset \{1,2,3\}$.

Linear Independence of a Set of Vectors

- The vectors $\{v_1, v_2, \dots, v_m\}$ are *linearly independent* if the *only* real numbers $\alpha_1, \alpha_2, \dots, \alpha_m$ yielding a linear combination of vectors that adds up to the zero vector,

$$\alpha_1 v_1 + \alpha_2 v_2 + \dots + \alpha_m v_m = 0_{n \times 1},$$

are $\alpha_1 = 0, \alpha_2 = 0, \dots, \alpha_m = 0$.

Suppose that $V \subset \mathbb{R}^n$ is a nonempty subset of \mathbb{R}^n .

Definition

V is a subspace of \mathbb{R}^n if any linear combination constructed from elements of V and scalars in \mathbb{R} is once again an element of V . One says that V is *closed under linear combinations*.

In symbols, $V \subset \mathbb{R}^n$ is a subspace of \mathbb{R}^n if for all real numbers α and β , and all vectors v_1 and v_2 in V

$$\boxed{\alpha v_1 + \beta v_2 \in V.}$$

Definition

Suppose that $S \subset \mathbb{R}^n$, then S is a set of vectors. The set of all possible linear combinations of elements of S is called the span of S ,

$\text{span}\{S\} := \{\text{all possible linear combinations of elements of } S\}.$

Suppose that V is a subspace of \mathbb{R}^n . Then $\{v_1, v_2, \dots, v_k\}$ is a *basis for V* if

- ▶ the set $\{v_1, v_2, \dots, v_k\}$ is linearly independent, and
- ▶ $\text{span}\{v_1, v_2, \dots, v_k\} = V$.
- ▶ The maximum number of vectors in any linearly independent set contained in V is the *dimension* of V (here k).

Definition

Let $n \geq 1$ and, as before, define $e_i := i$ -th column of the $n \times n$ identity matrix, I_n . Then

$$\{e_1, e_2, \dots, e_n\}$$

is a basis for the vector space \mathbb{R}^n .

Its elements e_i are called both natural (standard) basis vectors and canonical basis vectors.

We write a general system of linear equations as

$$a_{11}x_1 + a_{12}x_2 + \cdots + a_{1n}x_n = b_1$$

$$a_{21}x_1 + a_{22}x_2 + \cdots + a_{2n}x_n = b_2$$

$$\vdots = \vdots$$

$$a_{n1}x_1 + a_{n2}x_2 + \cdots + a_{nn}x_n = b_n$$

We can write this system as:

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix},$$

where:

$$A := \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix}, x := \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, b := \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}.$$

$$A = \begin{bmatrix} 3 & \mathbf{0} & \mathbf{0} \\ 2 & -1 & \mathbf{0} \\ 1 & -2 & 3 \end{bmatrix}$$

- ▶ All terms above the diagonal of the matrix A are zero.
- ▶ More precisely, the condition is $a_{ij} = 0$ for all $j > i$.
- ▶ Such matrices are called *lower-triangular*.

Lower Triangular Systems and Forward Substitution

We will solve this example using a method called *forward substitution*.

$$\begin{array}{rcl} 3x_1 & = & 6 \\ 2x_1 - x_2 & = & -2 \\ x_1 - 2x_2 + 3x_3 & = & 2 \end{array} \iff \underbrace{\begin{bmatrix} 3 & 0 & 0 \\ 2 & -1 & 0 \\ 1 & -2 & 3 \end{bmatrix}}_A \underbrace{\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}}_x = \underbrace{\begin{bmatrix} 6 \\ -2 \\ 2 \end{bmatrix}}_b.$$

$$A = \begin{bmatrix} 1 & 3 & 2 \\ \mathbf{0} & 2 & 1 \\ \mathbf{0} & \mathbf{0} & 3 \end{bmatrix}$$

- ▶ All terms below the diagonal of the matrix A are zero.
- ▶ More precisely, the condition is $a_{ij} = 0$ for $i > j$.
- ▶ Such matrices are called *upper-triangular*.

We solve the upper triangular systems using a method called *back substitution*.

$$\begin{array}{rcl} x_1 + 3x_2 + 2x_3 & = & 6 \\ 2x_2 + x_3 & = & -2 \\ 3x_3 & = & 4, \end{array} \iff \underbrace{\begin{bmatrix} 1 & 3 & 2 \\ 0 & 2 & 1 \\ 0 & 0 & 3 \end{bmatrix}}_A \underbrace{\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}}_x = \underbrace{\begin{bmatrix} 6 \\ -2 \\ 4 \end{bmatrix}}_b.$$

LU Factorization for Solving Linear Equations

- ▶ We wish to solve the system of linear equations $Ax = b$.
- ▶ If we can factor $A = L \cdot U$, where U is upper triangular and L is lower triangular. Then

$$L \cdot Ux = b.$$

- ▶ Define $U \cdot x =: y$, then

$$Ly = b$$

$$Ux = y.$$

- ▶ We first solve for y via forward substitution. Given y , we solve for x via back substitution.

Least Squares Solutions to $A_{n \times m} \cdot x_{m \times 1} = b_{n \times 1}$

- ▶ Assume $A^T A$ is invertible, i.e., the columns of A are linearly independent.
- ▶ Then there is a unique vector $x^* \in \mathbb{R}^m$ achieving $\min_{x \in \mathbb{R}^m} \|Ax - b\|^2$ and it satisfies the equation (called *the normal equations*)

$$(A^T A) x^* = A^T b.$$



$$x^* = (A^T A)^{-1} A^T b \iff x^* = \arg \min_{x \in \mathbb{R}^m} \|Ax - b\|^2 \iff (A^T A) x^* = A^T b.$$

Suppose that A is an $n \times m$ matrix with linearly independent columns.

Fact

Then there exists an $n \times m$ matrix Q with orthonormal columns, $Q^T Q = I$, and an upper triangular, $m \times m$, invertible matrix R such that $A = Q \cdot R$.

Least Squares via the QR Factorization

Whenever the columns of A are linearly independent, a least squared error solution to $Ax = b$ is computed as

- ▶ factor $A =: QR$,
- ▶ compute $\bar{b} := Q^T b$, and then
- ▶ solve $Rx = \bar{b}$ via back substitution.

- ▶ A function (or a map) view of a matrix defines two subspaces:
 - 1 its *null space* and
 - 2 its *range*.
- ▶ Let A be an $n \times m$ matrix.
- ▶ We can then define a function $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ by, for each $x \in \mathbb{R}^m$

$$f(x) := Ax \in \mathbb{R}^n.$$

The following subsets are naturally motivated by the function view of a matrix.

Definition

- 1 $\text{null}(A) := \{x \in \mathbb{R}^m \mid Ax = 0_{n \times 1}\}$ is the *null space* of A .
- 2 $\text{range}(A) := \{y \in \mathbb{R}^n \mid y = Ax \text{ for some } x \in \mathbb{R}^m\}$ is the *range* of A .

Range of A Equals Column Span of A

Let A be an $n \times m$ matrix, its columns are vectors in \mathbb{R}^n ,

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1m} \\ a_{21} & a_{22} & \cdots & a_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nm} \end{bmatrix} =: \begin{bmatrix} a_1^{\text{col}} & a_2^{\text{col}} & \cdots & a_m^{\text{col}} \end{bmatrix}$$

Then

$$\text{range}(A) := \{Ax \mid x \in \mathbb{R}^m\} = \text{span}\{a_1^{\text{col}}, a_2^{\text{col}}, \dots, a_m^{\text{col}}\} =: \text{col span}\{A\}.$$

Definition

For an $n \times m$ matrix A ,

1 $\text{rank}(A) := \dim \text{range}(A).$

2 $\text{nullity}(A) := \dim \text{null}(A).$

Because $\text{range}(A) \subset \mathbb{R}^n$, we see that $\text{rank}(A) \leq n$.

Theorem

For an $n \times m$ matrix A , we have the property

$$\text{rank}(A) + \text{nullity}(A) = m \quad \text{number of columns of } A.$$

- ▶ *Since $\text{rank}(A)$ is equal to the number of linearly independent columns of A , it follows that $\text{nullity}(A)$ is counting the number of linearly dependent columns of A .*
- ▶ *If all of the columns of A are linearly independent, then none are dependent, and hence $\text{null}(A) = \{0_{m \times 1}\}$.*

Multi-objective least squares

- ▶ goal: choose n -vector x so that k norm squared objectives

$$J_1 = \|A_1x - b_1\|^2, \dots, J_k = \|A_kx - b_k\|^2$$

are all small

- ▶ A_i is an $m_i \times n$ matrix, b_i is an m_i -vector, $i = 1, \dots, k$
- ▶ J_i are the objectives in a *multi-objective optimization problem* (also called a *multi-criterion problem*)
- ▶ could choose x to minimize any one J_i , but we want *one* x that makes them all small

Weighted sum objective

- ▶ choose positive *weights* $\lambda_1, \dots, \lambda_k$ and form *weighted sum objective*

$$J = \lambda_1 J_1 + \dots + \lambda_k J_k = \lambda_1 \|A_1 x - b_1\|^2 + \dots + \lambda_k \|A_k x - b_k\|^2$$

- ▶ we'll choose x to minimize J
- ▶ we can take $\lambda_1 = 1$, and call J_1 the *primary objective*
- ▶ interpretation of λ_i : how much we care about J_i being small, relative to primary objective
- ▶ for a bi-criterion problem, we will minimize

$$J_1 + \lambda J_2 = \|A_1 x - b_1\|^2 + \lambda \|A_2 x - b_2\|^2$$

Weighted sum minimization via stacking

- ▶ write weighted-sum objective as

$$J = \left\| \begin{bmatrix} \sqrt{\lambda_1}(A_1x - b_1) \\ \vdots \\ \sqrt{\lambda_k}(A_kx - b_k) \end{bmatrix} \right\|^2$$

- ▶ so we have $J = \|\tilde{A}x - \tilde{b}\|^2$, with

$$\tilde{A} = \begin{bmatrix} \sqrt{\lambda_1}A_1 \\ \vdots \\ \sqrt{\lambda_k}A_k \end{bmatrix}, \quad \tilde{b} = \begin{bmatrix} \sqrt{\lambda_1}b_1 \\ \vdots \\ \sqrt{\lambda_k}b_k \end{bmatrix}$$

- ▶ so we can minimize J using basic ('single-criterion') least squares

Weighted sum solution

- ▶ assuming columns of \tilde{A} are independent,

$$\begin{aligned}\hat{x} &= (\tilde{A}^T \tilde{A})^{-1} \tilde{A}^T \tilde{b} \\ &= (\lambda_1 A_1^T A_1 + \cdots + \lambda_k A_k^T A_k)^{-1} (\lambda_1 A_1^T b_1 + \cdots + \lambda_k A_k^T b_k)\end{aligned}$$

- ▶ can compute \hat{x} via QR factorization of \tilde{A}
- ▶ A_i can be wide, or have dependent columns

Optimal trade-off curve

- ▶ bi-criterion problem with objectives J_1, J_2
- ▶ let $\hat{x}(\lambda)$ be minimizer of $J_1 + \lambda J_2$
- ▶ called *Pareto optimal*: there is no point z that satisfies


$$J_1(z) < J_1(\hat{x}(\lambda)), \quad J_2(z) < J_2(\hat{x}(\lambda))$$

i.e., no other point x beats \hat{x} on both objectives



- ▶ *optimal trade-off curve*: $(J_1(\hat{x}(\lambda)), J_2(\hat{x}(\lambda)))$ for $\lambda > 0$

Example

A_1 and A_2 both 10×5



Objectives versus λ and optimal trade-off curve



Using multi-objective least squares

- ▶ identify the primary objective
 - the basic quantity we want to minimize
- ▶ choose one or more secondary objectives
 - quantities we'd also like to be small, if possible
 - *e.g.*, size of x , roughness of x , distance from some given point
- ▶ tweak/tune the weights until we like (or can tolerate) $\hat{x}(\lambda)$
- ▶ for bi-criterion problem with $J = J_1 + \lambda J_2$:
 - if J_2 is too big, increase λ
 - if J_1 is too big, decrease λ

Image de-blurring

- ▶ x is an image
- ▶ A is a blurring operator
- ▶ $y = Ax + v$ is a blurred, noisy image
- ▶ least squares de-blurring: choose x to minimize

$$\|Ax - y\|^2 + \lambda(\|D_v x\|^2 + \|D_h x\|^2)$$

D_v, D_h are vertical and horizontal differencing operations

- ▶ λ controls smoothing of de-blurred image

Example

blurred, noisy image



regularized inversion with $\lambda = 0.007$



Image credit: NASA

Regularization path

$$\lambda = 10^{-6}$$



$$\lambda = 10^{-4}$$



Regularization path

$$\lambda = 10^{-2}$$



$$\lambda = 1$$

